

The project will require a multiple linear regression model

In [318]:

```
# Importing Packages
!pip install statsmodels
import numpy as np
import pandas as pd
import seaborn as sns
import scipy.stats as stats
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
import statsmodels.formula.api as smf
import statsmodels.api as sm
from sklearn.metrics import r2_score
from sklearn.model_selection import KFold
```

```
Requirement already satisfied: statsmodels in
c:\users\msafw\anaconda3\lib\site-packages (0.14.4)
Requirement already satisfied: numpy<3,>=1.22.3 in
c:\users\msafw\anaconda3\lib\site-packages (from statsmodels) (1.24.2)
Requirement already satisfied: scipy!=1.9.2,>=1.8 in
c:\users\msafw\anaconda3\lib\site-packages (from statsmodels) (1.15.2)
Requirement already satisfied: pandas!=2.1.0,>=1.4 in
c:\users\msafw\anaconda3\lib\site-packages (from statsmodels) (2.2.3)
Requirement already satisfied: patsy>=0.5.6 in
c:\users\msafw\anaconda3\lib\site-packages (from statsmodels) (1.0.1)
Requirement already satisfied: packaging>=21.3 in
c:\users\msafw\anaconda3\lib\site-packages (from statsmodels) (24.2)
Requirement already satisfied: python-dateutil>=2.8.2 in
c:\users\msafw\anaconda3\lib\site-packages (from
pandas!=2.1.0,>=1.4->statsmodels) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in
c:\users\msafw\anaconda3\lib\site-packages (from
pandas!=2.1.0,>=1.4->statsmodels) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in
c:\users\msafw\anaconda3\lib\site-packages (from
pandas!=2.1.0,>=1.4->statsmodels) (2023.3)
Requirement already satisfied: six>=1.5 in
c:\users\msafw\anaconda3\lib\site-packages (from
python-dateutil>=2.8.2->pandas!=2.1.0,>=1.4->statsmodels) (1.17.0)
```

In [319]:

```
car = pd.read_csv("usedcars.csv")
car.head(5)
```

Out[319]:

Pric e	Ag e	Odomete r	Inspectio n	AB S	Sunroo f	Annualkm s
-----------	---------	--------------	----------------	---------	-------------	---------------

<b>0</b>	7.30	73	10.0	12	1	1	1.643836
<b>1</b>	3.85	115	30.0	20	1	0	3.130435
<b>2</b>	2.95	127	43.0	6	0	1	4.062992
<b>3</b>	4.80	104	54.0	25	1	1	6.230769
<b>4</b>	6.20	86	57.0	23	0	0	7.953488

In [320]:

```
car.var() # to check for scalability of the variables
```

Out[320]:

```
Price          1.541861
Age            362.853939
Odometer       2009.675408
Inspection     48.952391
ABS             0.215136
Sunroof        0.189336
Annualkms      21.515985
dtype: float64
```

In [321]:

```
car.shape
```

Out[321]:


```
(171, 7)
```

In [322]:

```
sns.histplot(car['Price'])
```

Out[322]:

```
<Axes: xlabel='Price', ylabel='Count'>
```




In [323]:

```
sns.histplot(car['Age'])
```

Out[323]:

```
<Axes: xlabel='Age', ylabel='Count'>
```




In [324]:

```
sns.histplot(car['Odometer'])
```

Out[324]:

```
<Axes: xlabel='Odometer', ylabel='Count'>
```




In [325]:

```
sns.histplot(car['Inspection'])
```

Out[325]:

```
<Axes: xlabel='Inspection', ylabel='Count'>
```




In [326]:

```
sns.histplot(car['Annualkms'])
```

Out[326]:

```
<Axes: xlabel='Annualkms', ylabel='Count'>
```



In [327]:

```
# Linear Regression without transformation of the variables
model = smf.ols('Price ~ Age + Odometer + Inspection + ABS + Sunroof +
Annualkms', data=car).fit()
print(model.summary())
```

```

OLS Regression Results

=====
=
Dep. Variable:          Price    R-squared:
0.641
Model:                  OLS      Adj. R-squared:
0.628
Method:                 Least Squares    F-statistic:
48.78
Date:                   Sun, 04 May 2025    Prob (F-statistic):
4.90e-34
Time:                   19:52:26    Log-Likelihood:
-191.59
No. Observations:      171    AIC:
397.2
Df Residuals:          164    BIC:
419.2
Df Model:               6
Covariance Type:       nonrobust
=====
=
```

	coef	std err	t	P> t	[0.025
0.975]					
-----					
Intercept	11.4162	1.044	10.931	0.000	9.354
Age	-0.0567	0.009	-6.066	0.000	-0.075
Odometer	0.0053	0.007	0.735	0.464	-0.009
Inspection	-0.0074	0.008	-0.878	0.381	-0.024
ABS	-0.2955	0.128	-2.314	0.022	-0.548
Sunroof	0.0358	0.136	0.263	0.793	-0.233
Annualkms	-0.1399	0.065	-2.142	0.034	-0.269
=====					
Omnibus:		4.617	Durbin-Watson:		
Prob(Omnibus):		0.099	Jarque-Bera (JB):		
Skew:		-0.006	Prob(JB):		
Kurtosis:		3.938	Cond. No.		
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.29e+03. This might indicate that there are strong multicollinearity or other numerical problems.

In [328]:

```
car['Age'] = np.log(car['Age'])
car['Inspection'] = np.log(car['Inspection'] +1)
car['Price'] = np.log(car['Price'])
car['Odometer'] = np.log(car['Odometer'])
car['Annualkms'] = np.log(car['Annualkms'])
# there is 1 value in the inspection of 0 which log does not work on so we
add the +1
```

In [329]:

```
car.corr()
```

Out[329]:




	Price	Age	Odometer	Inspection	ABS	Sunroof	Annualkm s
<b>Price</b>	1.000000	-0.707616	-0.568057	-0.044606	0.007920	-0.103568	-0.277357
<b>Age</b>	-0.707616	1.000000	0.395550	-0.010695	-0.047623	0.098166	-0.050854
<b>Odometer</b>	-0.568057	0.395550	1.000000	0.007959	-0.025866	0.123071	0.897141
<b>Inspection</b>	-0.044606	-0.010695	0.007959	1.000000	-0.111970	0.054090	0.013799
<b>ABS</b>	0.007920	-0.047623	-0.025866	-0.111970	1.000000	0.048747	-0.005221
<b>Sunroof</b>	-0.103568	0.098166	0.123071	0.054090	0.048747	1.000000	0.086611
<b>Annualkm s</b>	-0.277357	-0.050854	0.897141	0.013799	-0.005221	0.086611	1.000000

In [330]:

```
sns.boxplot(data = car)
```

Out[330]:

<Axes: >



In [331]:

```
car.head()
```

Out[331]:

	Price	Age	Odometer	Inspection	ABS	Sunroof	Annual kms
0	1.987874	4.290459	2.302585	2.564949	1	1	0.497032
1	1.348073	4.744932	3.401197	3.044522	1	0	1.141172
2	1.081805	4.844187	3.761200	1.945910	0	1	1.401920
3	1.568616	4.644391	3.988984	3.258097	1	1	1.829500

4 1.82454 4.45434 4.043051 3.178054 0 0 2.073611  
9 7

In [332]:

```
# Linear Regression transformations
model_transformation = smf.ols('Price ~ Age + Inspection + C(ABS) +
C(Sunroof) + Odometer', data=car).fit()
# C is statsmodel way of using a hot-encoding method in the model
print(model_transformation.summary())
```

```
OLS Regression Results

=====
=
Dep. Variable:          Price      R-squared:
0.603
Model:                OLS      Adj. R-squared:
0.591
Method:             Least Squares      F-statistic:
50.04
Date:                Sun, 04 May 2025      Prob (F-statistic):
2.35e-31
Time:                19:52:27      Log-Likelihood:
18.533
No. Observations:          171      AIC:
-25.07
Df Residuals:            165      BIC:
-6.215
Df Model:                5
Covariance Type:        nonrobust

=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
Intercept          7.8754      0.459      17.156      0.000      6.969
8.782
C(ABS) [T.1]       -0.0254      0.037      -0.687      0.493     -0.098
0.047
C(Sunroof) [T.1]   -0.0006      0.039      -0.014      0.989     -0.078
0.077
Age               -1.1030      0.103     -10.728      0.000     -1.306
-0.900
Inspection         -0.0337      0.032      -1.047      0.297     -0.097
0.030
Odometer           -0.2895      0.046      -6.355      0.000     -0.379
-0.200

=====
=
Omnibus:                2.695      Durbin-Watson:
2.302
```

Prob(Omnibus):	0.260	Jarque-Bera (JB):
2.245		
Skew:	-0.243	Prob (JB):
0.326		
Kurtosis:	3.280	Cond. No.
206.		

=====

=

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [333]:

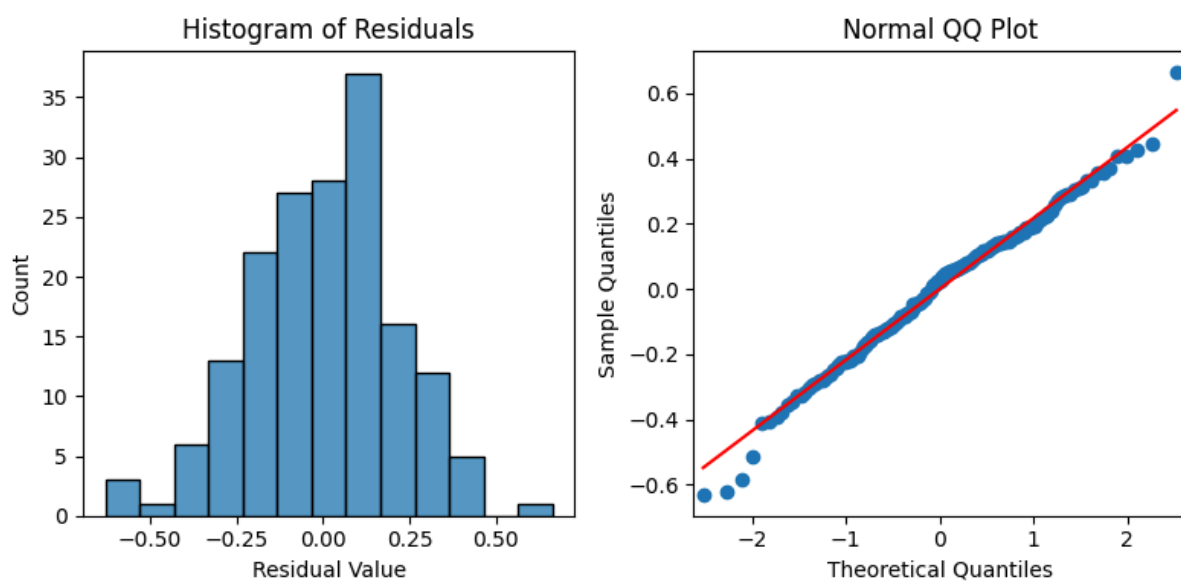
```
# Testing for Normality
# Plotting of Residuals
residuals = model_transformation.resid

# Histogram of Residuals
fig, axes = plt.subplots(1, 2, figsize = (8,4))
sns.histplot(residuals, ax=axes[0])
axes[0].set_xlabel("Residual Value")
axes[0].set_title("Histogram of Residuals")

# Q-Q plot of the residuals.
sm.qqplot(residuals, line='s',ax = axes[1])

axes[1].set_title("Normal QQ Plot")

plt.tight_layout()
# Show the plot.
plt.show()
```



In [334]:

```
# Testing independent observations
'''
Durbin-Watson: 2.3 is good score
The residuals are independent.
'''
```

Out[334]:

```
'\nDurbin-Watson: 2.3 is good score\nThe residuals are independent.\n'
```

In [335]:

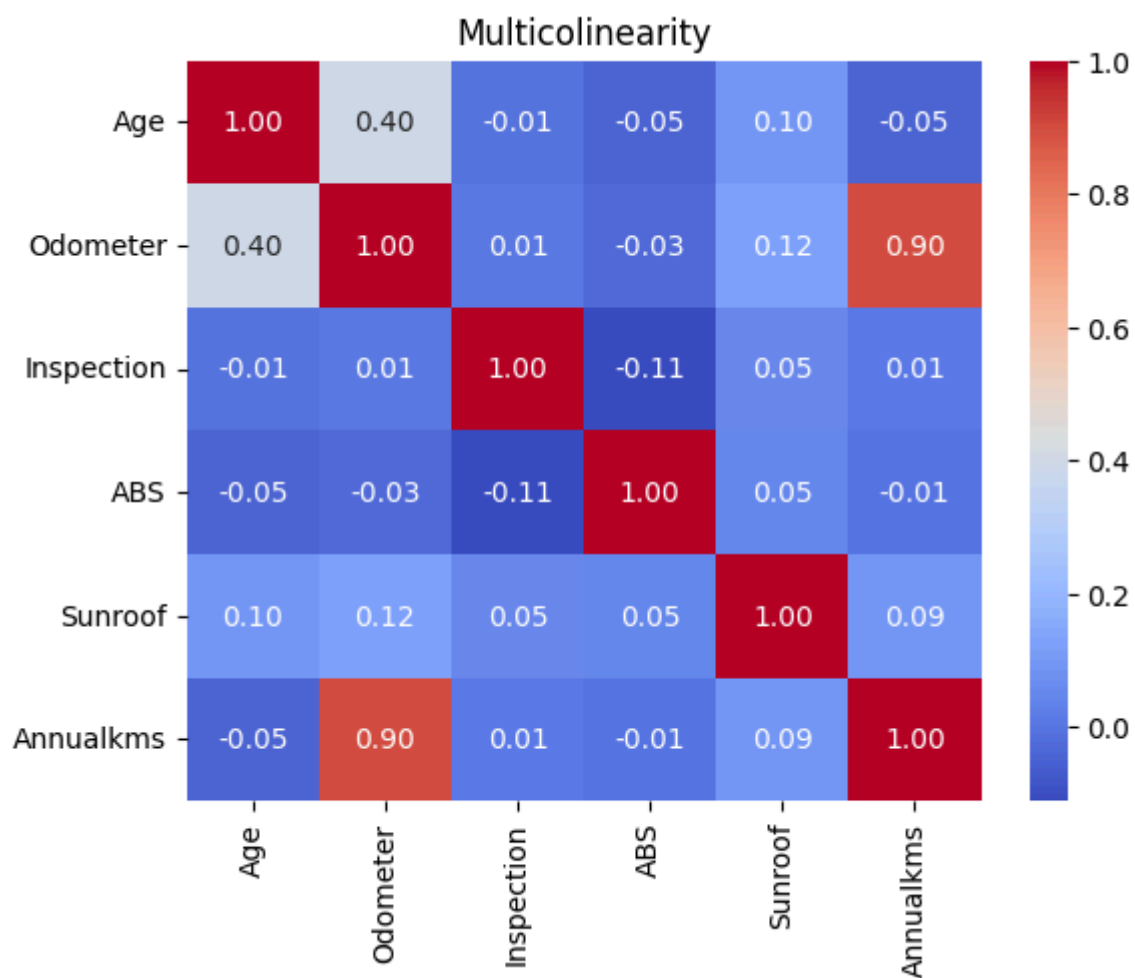
```
y = car['Price']
x = car.drop(columns='Price')
```

In [336]:

```
fig = sns.heatmap(x.corr(), annot=True, cmap="coolwarm", fmt=".2f")
fig.set_title("Multicollinearity")
```

Out[336]:

```
Text(0.5, 1.0, 'Multicollinearity')
```



In [337]:


```
# Checking for Homoscedasticity
```

```
fig = sns.scatterplot(x = model_transformation.fittedvalues, y =
model_transformation.resid)


fig.set_xlabel("Fitted Values")

fig.set_ylabel("Residuals")

fig.set_title("Fitted Values v. Residuals")
fig.axhline(0)
plt.show()
```



```
In [346]:
sm.graphics.influence_plot(model_transformation, criterion="cooks", size=8)
plt.title("Leverage vs. Residuals Squared (Influence Plot)")
plt.tight_layout()
plt.show()
```



In [339]:

```
rows_to_drop_of_car = [24,62,78] # Delete extreme outliers
car = car.drop(rows_to_drop_of_car, axis=0)
car.shape
```

Out[339]:

(168, 7)

In [340]:

```
# Linear Regression transformation of the variable age and inspection
model_transformation = smf.ols('Price ~ Age + Inspection + C(ABS) +
C(Sunroof) + Annualkms', data=car).fit()
# C is statsmodel way of using a hot-encoding method in the model
print(model_transformation.summary())
```

```

=====
                        OLS Regression Results
=====
=
Dep. Variable:          Price    R-squared:
0.645
Model:                  OLS      Adj. R-squared:
0.634
Method:                 Least Squares    F-statistic:
58.96
Date:                   Sun, 04 May 2025    Prob (F-statistic):
1.00e-34
```

```

Time: 19:52:30 Log-Likelihood:
29.751
No. Observations: 168 AIC:
-47.50
Df Residuals: 162 BIC:
-28.76
Df Model: 5
Covariance Type: nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025
Intercept	8.7678	0.451	19.420	0.000	7.876
C(ABS) [T.1]	-0.0282	0.035	-0.812	0.418	-0.097
C(Sunroof) [T.1]	0.0162	0.037	0.438	0.662	-0.057
Age	-1.4256	0.089	-15.940	0.000	-1.602
Inspection	-0.0323	0.030	-1.070	0.286	-0.092
Annualkms	-0.2974	0.043	-6.977	0.000	-0.382

```

=====
=
Omnibus: 0.013 Durbin-Watson:
2.123
Prob(Omnibus): 0.994 Jarque-Bera (JB):
0.110
Skew: 0.004 Prob(JB):
0.946
Kurtosis: 2.875 Cond. No.
180.
=====
=

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [341]:

```

new_car = pd.DataFrame({
    'Age': np.log(124),
    'Annualkms': np.log(12),
    'Inspection': np.log(9),
    'Sunroof': [1],
    'ABS': [0]
})

predicted_price = model_transformation.predict(new_car)

```



```
print(f"Predicted price: {np.exp(predicted_price[0]):.3f} euros")
```

```
Predicted price: 3.012 euros
```

In [342]:

```
new_car = pd.DataFrame({
    'Age': np.log(148),
    'Annualkms': np.log(40),
    'Inspection': np.log(9),
    'Sunroof': [1],
    'ABS': [0]
})
```

```
predicted_price = model_transformation.predict(new_car)
print(f"Predicted price: {np.exp(predicted_price[0]):.3f} euros")
```

```
Predicted price: 1.636 euros
```

In [343]:

```
new_car = pd.DataFrame({
    'Age': np.log(148),
    'Annualkms': np.log(40),
    'Inspection': np.log(9),
    'Sunroof': [1],
    'ABS': [1]
})
```

```
predicted_price = model_transformation.predict(new_car)
print(f"Predicted price: {np.exp(predicted_price[0]):.3f} euros")
```

```
Predicted price: 1.590 euros
```

In [344]:

```
# Train-test car database
```

```
train_data, test_data = train_test_split(car, test_size=0.2, random_state=42)
```

```
model_transformation = smf.ols('Price ~ Age + Inspection + C(ABS) +
C(Sunroof) + Annualkms', data=train_data).fit()
# you would normally fit and transform the training data but log must be
applied for both fitting and transforming, as we want the same scale for
both
```

```
y_pred = model_transformation.predict(test_data)
```

```
r2_test = r2_score(test_data['Price'], y_pred)
```

```
print("Train R-squared:", model_transformation.rsquared)
print("Test R-squared:", r2_test)
```

```
Train R-squared: 0.6550245493686415
```

Test R-squared: 0.5911407009740445

In [345]:

```
# Cross-validation
kf = KFold(n_splits=10, shuffle=True, random_state=42)
cv = []

for train_index, val_index in kf.split(train_data):
    fold_train = train_data.iloc[train_index] # for training
    fold_validate = train_data.iloc[val_index] # for validating

    model_transformation = smf.ols('Price ~ Age + Inspection + C(ABS) +
C(Sunroof) + Annualkms', data=fold_train).fit()
    # It did take a while to understand that looping into fold-train than
    actual train data, using k-fold
    y_val_pred = model_transformation.predict(fold_validate)
    # You do predict the scaled data by log for the validating
    r2 = r2_score(fold_validate['Price'], y_val_pred)
    cv.append(r2)

print("Cross-validated R-squareds result:", cv)
print(np.quantile(cv, [0.025, 0.975]))
print("Average CV R-squared result:", np.mean(cv))
```

Cross-validated R-squareds result: [0.8511238289233365, 0.7110102955455111,  
0.5661451178109761, 0.6207352257491261, 0.4161020626098847,  
0.33504571690299334, 0.2984171838379017, 0.6914932462306971,  
0.43347470692278967, 0.7569065789211618]  
[0.3066586 0.82992495]

Average CV R-squared result: 0.5680453963454377